

AI Speech Platform Buyability Benchmark

A benchmark of how AI Speech companies equip developers and operators to evaluate, budget for, and deploy with confidence.

12 companies · 3 sub-types · 8 dimensions

By Michelle Perkins
Founder, ValueTempo

Contents

HOW TO USE THIS REPORT

Founders: benchmark your commercial transparency against category leaders.
GTM operators: know exactly where each vendor's evidence holds and where it doesn't. Platform builders: map production risk before you commit.

BENCHMARK CONTEXT

- 03 Executive Summary
 - 04 What the Benchmark Measures
 - 05 Benchmark Overview
-

SIX KEY FINDINGS

- 06 Finding 1: Developer-First GTM Sets a New Transparency Floor
 - 07 Finding 2: Four Dimensions Scored Perfectly — Two Exposed the Shared Gap
 - 08 Finding 3: Vague Performance Claims Are a GTM Liability
 - 09 Finding 4: Nine of 12 Companies Leave Production Risk Undocumented
 - 10 Finding 5: Sub-Segment Comparison — Who Leads and Why
 - 11 Finding 6: A Perfect Score Is Not the Ceiling
-

NEXT STEPS

- 12 Operator Takeaways
 - 13 Methodology Notes
 - 14 Request a Score Walkthrough
-

APPENDIX

- 15 More in the June Benchmark
- 16 Also in the AVS Benchmark Series
- 17 Companies Evaluated

Executive Summary

The June 2026 AI Speech Platform Buyability Benchmark scored 12 AI Speech companies — spanning speech-to-text (STT) APIs, text-to-speech (TTS) platforms, and voice agent orchestration layers — on 8 dimensions of publicly available commercial evidence. The score measures what a buyer can independently verify before engaging sales. It does not measure product quality, brand strength, customer satisfaction, or market leadership.



What stood out

1. **86% category average — 17 points above May's top category.** Developer-first, usage-based GTM creates a structural transparency advantage.
2. **All 12 companies scored Exemplary — a benchmark first.** Four of eight dimensions scored perfectly across the board. The question has shifted from whether buyers can get started to what gaps remain at the top.
3. **Eight of 12 companies publish performance claims buyers cannot independently verify** — extending the sales cycle by forcing proof-of-concept evaluations that a published benchmark could eliminate.
4. **Nine of 12 companies do not document what happens when a deployment hits a rate limit or exhausts credits** — leaving production risk unresolved for every operator building on this infrastructure.
5. **Deepgram and Retell AI both scored 16/16 and still have named gaps in spend controls, billing disputes, and refund policies that informed buyers will surface.**

BOTTOM LINE:

AI Speech has solved for pricing transparency. The next frontier is performance verifiability, production risk documentation, and post-purchase commercial terms.

What the Benchmark Measures

The AI Speech Platform Buyability Benchmark measures the published commercial evidence developers and operators need to evaluate, budget for, and deploy with confidence.

What we measured

We reviewed publicly available commercial evidence across 12 AI Speech companies and 8 scoring dimensions.

The benchmark evaluates whether developers and operators can independently verify value, cost structure, and production risk from public sources alone.

Core question:

Can a buyer independently understand what the product does, who it is for, what it costs, what they are paying for, what drives usage, what happens at limits, and what trust signals are available?

What we did not measure

The score does not measure product quality, brand strength, market leadership, customer satisfaction, revenue performance, security posture, or long-term company credibility.

A higher score indicates stronger published commercial evidence.

Benchmark Overview

All 12 companies scored Exemplary — the first category in the benchmark where this has happened.

The AI Speech market has reached a transparency baseline that most GTM software categories have not. Usage-based pricing, developer-first distribution, and API-native go-to-market have collectively produced a category where every company meets the standard for buyer-ready commercial evidence.

AI Speech Platform AVS Scores — June 2026




Figure 1 — Score distribution: 12 companies, ranked highest to lowest, color-coded by sub-type. Scores reflect publicly observable commercial evidence as of June 30, 2026. N=12.

What this means

All 12 companies achieved the **Exemplary** threshold — a benchmark first. The cluster above 80% signals category maturity. The 19-point gap between the top pair (100%) and the lowest-scoring companies (81%) reflects documentation depth rather than product capability.

Finding 1: Developer-First GTM Sets a New Transparency Floor

Category averages hide where the market is actually moving. AI Speech has moved furthest.




Figure 2 — Category average comparison: AI Speech Platform (86%) vs. May 2026 categories (48%–69%).

What this means

The 86% average is 17 points above May's AI SaaS Buyability Benchmark. In AI speech, the developer evaluates before talking to anyone — the pricing page, documentation, and trust surfaces carry the buyer education that a discovery call handles in sales-led categories. That's why the four dimensions usage-based pricing solves by default all scored perfectly here, while they remain the hardest dimensions across May's five categories.

The gaps that remain require deliberate investment regardless of pricing model. Consumption companies averaged 92.5% vs. 83.75% for hybrid — but the gap concentrates in two dimensions: verifiable performance claims (D1) and production-edge behavior (D7). A single value unit is easier to anchor a benchmark claim to, and a single exhaustion event is easier to document than the two-layer failure modes of a hybrid model.

Finding 2: Four Dimensions Scored Perfectly. Two Exposed the Shared Gap.

The strongest and weakest dimensions reveal where the category has arrived and where it still needs to go.

Dimension Score Profile — AI Speech Platform, June 2026



Figure 3 — Each axis = average score out of 2.0 across 12 companies.
Flat edges indicate gap dimensions.

What this means

The polygon's near-perfect shape on D2, D3, D4, and D6 shows category-wide consensus on how AI Speech companies present themselves. The flat edges at D1 (avg 1.33) and D7 (avg 1.25) are consistent gaps across the category. Nine of twelve companies left overages unclear.

Finding 3: Vague Performance Claims Are a GTM Liability

Eight of 12 companies make claims buyers cannot verify. That forces a proof of concept where a benchmark citation would do.

D1 — Product North Star: Verifiable Claims vs. Unverifiable Claims




Figure 4 — D1 score distribution: companies by product north star score

What this means

Eight of 12 companies publish performance claims — "ultra-low latency," "human-like voices," "state-of-the-art accuracy" — without anchoring them to a number, external benchmark, or independent evaluation. Serious buyers respond by running a proof of concept, turning a documentation gap into weeks of engineering time on both sides. Deepgram, Retell AI, Rime, and Speechmatics published verifiable evidence and eliminated that cost from the evaluation.

Finding 4: Nine of 12 Companies Leave Production Risk Undocumented

The shared floor of this category is the missing documentation of what happens when production hits a limit.

D7 — Overages & Risk Allocation: Documented vs. Undocumented Edge Behavior




Figure 5 — D7 measures whether a company documents what happens when usage limits or credits are exhausted.

What this means

Nine of 12 companies document that rate limits exist but not what happens when they're hit — "whether an API call fails gracefully or silently, whether a notification fires, what happens when usage limits are reached mid-call. For operators building production systems, that is an unquantified risk. Deepgram, Retell AI, and LMNT documented the edge cases; the remaining nine have not.

Finding 5: Sub-Segment Comparison: Who Leads and Why

Full-stack platforms and voice agent builders outperform pure-play TTS on average. The reason is structural.

Average AVS Score by Sub-type — AI Speech Platform, June 2026




Figure 6 — Sub-segment average bar chart: Full Speech Platform, Voice Agent Platform, TTS, STT.

What this means

Full Speech Platform scores highest at 93.8%, reflecting the documentation advantage of a unified STT + TTS pricing framework. Voice agent platforms follow at 89.6% — billing across LLM tokens, STT minutes, TTS characters, and PSTN pass-through forces documentation discipline on each component. TTS specialists face a higher bar on D8: voice cloning triggers consent flow requirements that pure STT companies are not subject to. ElevenLabs, Murf AI, Rime, and WellSaid Labs published explicit consent flows and score D8=2. Cartesia and LMNT have not.

Finding 6: A Perfect Score Is Not the Ceiling

Two companies scored 100%. Both have specific, named gaps that informed buyers will find.

What this means

Deepgram and Retell AI both earned perfect scores on all eight dimensions. But both companies still have gaps that live above the rubric — in a layer of commercial terms, in-product controls, and post-purchase trust signals that no dimension currently captures:

Deepgram:

- Auto-reload triggers at a fixed threshold (\$100 when balance hits \$10); no configurable spend cap, adjustable alert, or usage monitoring dashboard is published
- Refund policy covers unused credits within 30 days; no dispute process exists for contested charges or platform-error-induced billing
- Pricing calculator covers TTS concurrency only; no unified estimator for mixed STT + TTS + Voice Agent workloads

Retell AI:

- No refund or credit-back policy documented for any usage scenario
- Payment failure handling (suspension timeline, grace period) is not documented on any buyer-facing public page.
- No spend-spike alerting or anomaly detection documented on any public page

BOTTOM LINE:

Deepgram and Retell AI clear every rubric dimension. The gaps that remain sit at the edge of what the current D7 threshold requires: billing disputes, usage spikes, refund eligibility, and spend controls. That is the next layer this category needs to solve.

Operator Takeaways

The category has solved for pricing transparency. The next advantage comes from verifiability and production-readiness documentation.

1

Publish one verifiable performance number

Name a benchmark, cite a position on TTS-Arena or Artificial Analysis, or publish a WER result. A buyer who can verify a performance claim before the sale does not need a proof of concept to establish it.

2

Document what happens at production limits

Publish what error code fires at rate limit, whether auto-top-up is opt-in or opt-out, and what happens to in-flight requests at exhaustion. For voice agent platforms, document what happens when usage limits are reached — call termination behavior, notification timing, and billing edge cases.

3

Match cost driver documentation to your sub-type

STT buyers need to know how model tier, language, and add-ons affect cost. TTS buyers need concurrent stream limits and voice premiums. Voice agent buyers need the full component stack.

4

Invest in commercial terms completeness beyond the score

Spend caps, billing dispute processes, and refund policies are not rubric dimensions — but they are questions enterprise operators ask before signing.

Methodology Notes

How to read the score

The AVS Rubric evaluates publicly available buyer-facing evidence: the same information a buyer can review before engaging sales.

Each company was scored across 8 dimensions, with 0 to 2 points per dimension and 16 total points available.

0 = Evidence is missing or not independently usable

1 = Evidence is partially available, but incomplete

2 = Evidence is specific enough to support independent buyer evaluation

Category-specific scoring guide

AI Speech Platform introduces sub-type interpretation rules for three product types: STT APIs, TTS platforms, and voice agent orchestration layers. Rules cover voice cloning consent thresholds (D8), PSTN pass-through documentation requirements (D5), credit conversion requirements (D4), and API rate limit vs. mid-call behavior differentiation (D7). Full scoring guide available in the analyst report.

What the score reflects

A higher score means stronger published commercial evidence. It should not be read as a proxy for product quality, company strength, market position, or customer satisfaction.

Point-in-time snapshot

Scores reflect publicly available evidence at the time of review (June 30, 2026).

Company websites, pricing pages, documentation, and packaging may change after publication. No company paid for inclusion or received advance notice of scores.

Request a Score Walkthrough

The AVS Benchmark is designed to help AI SaaS teams understand how their public commercial evidence supports or slows independent buyer evaluation.

A score walkthrough helps identify:

1 Where your evidence is already strong

The dimensions where buyers can understand value, cost, risk, and readiness without extra explanation.

2 Where buyers still need sales to fill gaps

The missing or incomplete evidence that makes evaluation harder before the first call.

3 What to publish next

The highest-leverage changes to improve buyer confidence, reduce commercial friction, and make the product easier to evaluate.

**REQUEST A
WALKTHROUGH**

Schedule a free walkthrough →
calendly.com/mlhperkins/30min

More in the June Benchmark

This public report covers the six key findings and category-level patterns. The full analyst report goes significantly deeper:

1 Dimension-level scores for all 12 companies

The exact breakdown of which dimensions each company scored 1 vs. 2 on, and the specific evidence that determined each score

2 12 individual company profiles

What each company does better than category peers, named improvement areas per dimension, and the specific documentation changes that would move each score

3 How the scoring guide rules were applied

The PSTN pass-through rule, voice cloning consent threshold, credit conversion rule, and sub-type differentiation (API rate limits vs. mid-call behavior) applied company by company

4 Sub-segment analysis by dimension

Which dimensions differentiate STT vs. TTS vs. voice agent platforms within the Exemplary band, and what that reveals about how each sub-type approaches buyer transparency

5 Cost scenario framework

Detailed worked examples for STT, TTS, and voice agent use cases using published rate cards, showing what self-serve cost estimation looks like for each sub-type

**REQUEST A
WALKTHROUGH**

Schedule a free walkthrough →
calendly.com/mlhperkins/30min

This report is published by ValueTempo for informational and educational purposes. The AVS Benchmark evaluates publicly available commercial evidence published by AI SaaS companies at the time of review. Scores reflect the quality, specificity, and accessibility of buyer-facing evidence, not product quality,

company performance, brand strength, customer satisfaction, security posture, or market leadership. Company names, product names, logos, and trademarks referenced in this report are the property of their respective owners and are used for identification and analysis only.

Also in the AVS Benchmark Series

May 2026 AI SaaS Buyability Benchmark

The benchmark that established buyability as a measurable GTM standard. 60 companies across 5 categories — **Sales Intelligence, AI Sales Engagement, Revenue Intelligence, Agent Platform, and Conversation Intelligence**. Category averages ranged from 58% to 69%.

Key finding: Most companies publish enough evidence to be considered. Making independent evaluation easy requires more. The value unit is the single strongest separator.

The May and June benchmarks use the same 8-dimension AVS Rubric, making scores directly comparable across categories and over time.

Download the May 2026 benchmark:

valuetempo.com/ai-saas-buyability-benchmark-may-2026

This report is published by ValueTempo for informational and educational purposes. The AVS Benchmark evaluates publicly available commercial evidence published by AI SaaS companies at the time of review. Scores reflect the quality, specificity, and accessibility of buyer-facing evidence, not product quality,

company performance, brand strength, customer satisfaction, security posture, or market leadership. Company names, product names, logos, and trademarks referenced in this report are the property of their respective owners and are used for identification and analysis only.

Appendix: Companies Evaluated

AI SPEECH PLATFORM (12) — June 2026

COMPANY	SUB-TYPE	PRICING MODEL*	BAND
AssemblyAI	STT	Consumption	Exemplary
Bland.ai	Voice Agent Platform	Hybrid	Exemplary
Cartesia	TTS / STT / Voice Agent	Hybrid	Exemplary
Deepgram	Full Speech Platform	Consumption	Exemplary
ElevenLabs	TTS	Hybrid	Exemplary
LMNT	TTS	Hybrid	Exemplary
Murf AI	TTS — Creative voice	Hybrid	Exemplary
Retell AI	Voice Agent Platform	Consumption	Exemplary
Rime	TTS	Consumption	Exemplary
Speechmatics	Full Speech Platform	Consumption	Exemplary
Vapi	Voice Agent Platform	Hybrid	Exemplary
Wellsaid Labs	TTS — Enterprise	Hybrid	Exemplary

* Pricing model mix: 7 Hybrid (58%) · 5 Consumption (42%).

Band definitions: Developing, Credible, Trusted, and Exemplary reflect the strength of published commercial evidence at the time of review. Bands do not measure product quality, company performance, customer satisfaction, security posture, or market leadership.